

MASTER THESIS
COMPUTER SCIENCE

Scaling UIMA with Apache Spark

Simon Gehring

Am Jesuitenhof 3
53117 Bonn
gehring@uni-bonn.de
Matriculation Number 2553262

At the
RHEINISCHE FRIEDRICH-WILHELMS-UNIVERSITÄT BONN

supervised by
Prof. Dr. Heiko RÖGLIN and Dr. Timm HEUSS

September 21, 2018

Contents

Contents	i
1 Introduction	1
1.1 Motivation	1
1.2 Implementation Specification	2
1.2.1 CPU & RAM Usage	2
1.2.2 Throughput	3
1.2.3 Maintainability	3
1.2.4 Implementability	3
1.2.5 Code Quality	3
1.3 Outline	3
2 Basics	5
2.1 UIMA-Family	5
2.1.1 Apache UIMA	6
2.1.2 UIMAFit	9
2.1.3 UIMA-CPE	9
2.1.4 UIMA-AS	11
2.2 Distributed Computation	12
2.2.1 MapReduce	13
2.2.2 Resilient Distributed Datasets	15
2.3 Containerization and Docker	16
3 Scaling UIMA with Spark	18
3.1 Selection of Technologies	18
3.2 Implementation	18
3.2.1 Initialization	19
3.2.2 Transport	19
3.2.3 Process	22
3.2.4 Result	22
3.3 Data Distribution	23
3.3.1 Serialization	23
3.3.2 Compression	24

4	Evaluation	28
4.1	Conceptual Comparison	28
4.2	Setup	29
4.2.1	Hard- and Software	29
4.2.2	Apache Spark	29
4.2.3	UIMA-AS	31
4.2.4	Analysis Engines	32
4.2.5	Document Corpus	33
4.3	Results	33
4.3.1	CPU Usage	33
4.3.2	RAM Usage	33
4.3.3	Document Throughput	34
4.3.4	Byte Throughput	35
4.3.5	Maintainability	37
4.3.6	Implementability	38
4.3.7	Code Quality	39
5	Summary	40
5.1	Limitations	40
5.2	Related Work	40
5.2.1	Leo	41
5.2.2	v3NLP	42
5.3	Conclusion	42
5.4	Availability	43
5.5	Outlook	43
	Bibliography	I

Abstract

In this thesis, we will introduce a new method of scaling UIMA (Unstructured Information Management Architecture) by distributing its workload to different machines inside a Spark (Apache Spark) cluster. UIMA, being a general purpose NLP (Natural Language Processing) framework, already ships with various methods of scaling. With one method being UIMA-CPE (UIMA Collection Processing Architecture), an outdated framework and another being UIMA-AS (UIMA Asynchronous Scaleout), a newer, service-based idea, expecting the developer to many XML (Extensible Markup Language) files for every configuration detail of the deployment, a new approach seems necessary.

Given an instance of a Spark cluster, the framework presented in this thesis will be easy to implement and maintain. Simultaneously, it will pose almost no restrictions to the developer in terms of functionality of UIMA. This especially includes the reusing of already established Analysis Engines, as for example contained in the large DKPro Core (Darmstadt Knowledge Processing Software Repository).

After introducing the framework itself and going into detail of its functionalities, we will evaluate it against UIMA-AS. For this, we will introduce metrics on how to measure the quality of either framework. We will also discuss maintainability and implementability of the framework and compare its code quality with the one of UIMA-AS.

At the end we reach the conclusion that both approaches are Pareto efficient, having use cases for each of the framework where it excels the other.

Chapter 1

Introduction

Natural language is most commonly used to transmit information human-to-human. While most of this interaction takes place orally or written on paper, the digital revolution and the rise of social media increased the amount of digitally stored natural language tremendously. Many opportunities arise from this amount of digital data, specifically in the field of machine learning, semantic and intelligent systems. In 2011, IBM’s QA (Question Answering) system “Watson” famously outmatched professional players in the quiz show “Jeopardy!” [Fer12, ESI⁺12] and is now used in the *Crew Interactive Mobile Companion* (CIMON), an AI aiding with various tasks in space [NAS]. Kudesia et al. proposed 2012 an algorithm to detect so called CAUTIs¹, common hospital-acquired infections, by utilizing a NLP analysis with precomputed language models on the medical records of patients [KSDG12].

Current projections suggest a growth of revenues from the natural language processing market in North America, Western Europe, and the Asia-Pacific region of at least a factor of three until 2024 [Stad, Stab, Stac].

1.1 Motivation

Natural language is inherently unstructured and hardly machine readable. Even a seemingly easy task like separating a sentence into words is still an ongoing research topic [PT18]. Since many NLP frameworks like Stanford’s Core NLP Suite [MSB⁺14], The Natural Language Toolkit [BL04] and Apache OpenNLP [Apaa] exist, there was a need for generalizing the NLP approach. In 1995, the University of Sheffield began developing GATE (General Architecture for Text Engineering), a graphical development application for generic NLP problems [CMBT02]. Aside from algorithms for typical NLP tasks like tokenization, sentence splitting or part of speech tagging, GATE provides a graphical interface for developing custom algorithms in form of plug-ins, which can use the results of previous NLP analyses. However, the possibilities for scaling GATE applications are sparse. In 2012 GATECloud.net launched, a proprietary, cloud based GATE computa-

¹Catheter-associated Urinary Tract Infections

tion interface [TRCB13]. Since then, no improvements to the scaling capability of GATE is publicly known.

In [FL04], Ferrucci and Lally introduced UIMA, another general purpose NLP framework. Initially developed by IBM, UIMA has been open-source and is maintained by Apache since 2006. UIMA itself provides no built-in NLP analyses, but offers a common analysis structure among its plug-ins, which is used to combine different NLP approaches into one single framework, just like GATE. A popular implementation example of UIMA is IBM’s QA system “Watson” [Fer12, ESI⁺12], as described above. For this matter, UIMA was configured to run on thousands of processor cores to achieve a feasible reaction time [ESI⁺12]. Although this example seems to demonstrate that Apache UIMA is indeed very scalable, it has its drawbacks. The native UIMA scaling framework UIMA-AS is configured by XML files, which are difficult to maintain. Furthermore, it does not work out-of-the-box with other UIMA derivatives like UIMAFit (Factories, Injection, and Testing library for UIMA).

In this thesis a scale-out framework for UIMA, which utilizes modern technology to ensure maintainability and scalability will be implemented and evaluated.

1.2 Implementation Specification

The implementation should meet a number of requirements. First and foremost, the underlying framework, Apache UIMA, must not be limited in functionality. Since one of UIMAs strengths is the modularity and ease of plug-in development, the scale-out framework to implement is required to work with native UIMA classes without any restriction. While this requirement sounds easy to meet, it is actually quite limiting since UIMA plug-ins can be arbitrary Java code. Without special care, such code is not necessarily thread-safe and thus can not be safely executed multiple times in parallel within one JVM (Java Virtual Machine). Especially code that already exists for UIMA must be able to be reused within the framework.

There are a number of metrics the framework should optimize. The last three metrics maintainability, implementability and code quality, are not trivial to measure, since those are subject to individual perception. However, implementability can be measured in LoC (Lines of Code) needed to achieve a common use case like deploying a predefined NLP algorithm. There are further static code analyzers, which return an amount of code quality issues. Such a score can be used to measure the quality of the framework’s code. This leaves maintainability as the only attribute that can not easily be measured in numbers. However, large differences between the different approaches become clear in Chapter 4, making a comparison possible.

1.2.1 CPU & RAM Usage

With CPU and RAM Usage being the proportion of actually used (virtual or real) resources, higher values are more preferable. Even if many of those resources go into administrative tasks, a scaling framework should use as many of the available resources

as possible, when faced with a large list of parallel jobs. Thus, a value near one should be aimed for. At the same time, resource allocation when idling should be as low as possible.

1.2.2 Throughput

A little more obvious metric is the achieved throughput of data. Since collections of input documents can be shaped very differently, both, a high document and byte throughput are desired. When handling small documents, maybe even single words or sentences, the byte throughput is limited by the actual input size and the large number of documents is responsible for the size of the input data collection. In such a situation, a high document throughput would be preferred to a high byte throughput. On the other hand, on larger documents, a high byte throughput is the more accurate metric, since it is independent of the individual size of the input documents.

1.2.3 Maintainability

The framework is aimed to work in large academic but also industry compliant environments. Therefore, the maintainability is important. It describes the amount of effort to maintain any current usage of the framework, for example changing the underlying NLP algorithm, modifying the hardware setup or making configuration changes. This is inherently more complicated by the genericness of UIMA, which allows for sophisticated plug-in initialization and configuration logic, making on-the-fly changes more difficult.

1.2.4 Implementability

Furthermore, the framework should be easy to utilize. Code that already has been written for single-threaded execution should be easily reusable. This is especially important for UIMA since large repositories of plug-ins like the DKPro Core [DKPb] already exist and are infeasible to be rewritten.

1.2.5 Code Quality

Equally important as ease of utilization is the longevity of the framework. Ensuring a maximum of Code Quality is necessary to avoid large refactorings and API (Application Programming Interface) changes in the near future. In a non-academic environment, applications often have to last for a long time before replacement. Thus, a robust underlying code is desirable.

1.3 Outline

First, in Chapter 2, the functionality of UIMA is detailed, especially with focus on distributed computing and scaling. For this matter, both native UIMA scaling frameworks, UIMA-AS and CPE (Collection Processing Engine) are subject in said chapter. Also

Spark, a cluster-computing framework, will be briefly explained since it will form the foundation of the frameworks scaling capabilities. Docker follows, which is heavily used in the evaluation as virtualization technique.

Chapter 3 will explain the scaling framework in detail. It starts with Spark as the chosen technology and then proceeds to the implementation. A special focus lies on the data distribution between different worker machines.

Then, Chapter 4 deals with the results of the framework implementation. For this, the requirements given in Section 1.2 are evaluated against the framework, UIMA-AS and the single-threaded approach. A sophisticated setup, which is also depicted in the chapter, was necessary to provide comparable result.

Lastly, Chapter 5 summarizes the results, including possible limitations of the implementation. It also gives an overview over related work and an outlook on how to extend and publicize the framework.

Chapter 2

Basics

In this chapter, the two most important technologies for the framework are explained. This is necessary to get an understanding of the technical difficulties it faces and how the underlying concepts work. First, in Section 2.1 UIMA is introduced. After an in-depth introduction into the framework originally designed by IBM, UIMAFit will be explained. UIMAFit builds on top of UIMA, providing the developer with a native Java interface for creating and instantiating plug-ins. Related to the framework introduced in Chapter 3 are the two native scaling frameworks UIMA-CPE and UIMA-AS.

The second section of this chapter will be about Spark. While no advanced knowledge is needed to comprehend the usage of Spark as a distributed computation framework, it will still be a substantial part of the UIMA scaling framework in Chapter 3. Thus, an overview of its structure and distribution algorithm will be given. Although it was not used in the development of the framework, Docker will be introduced, since it was heavily utilized in the evaluation described in Chapter 4.

Although most important aspects and concepts of UIMA are also defined in the specifications, some minor changes and additions were made in the implementations. Since the framework must handle the actual implementation, all the presented concepts will be taken from Apache UIMA instead of the UIMA specification of 2009.

2.1 UIMA-Family

Unstructured Information Management Architecture (UIMA) Version 1.0 is an OASIS (Organization for the Advancement of Structured Information Standards) standard from 2009¹ that defines an interface for software components, or plug-ins, which are called analytics. Those analytics are supposed to analyze unstructured information and assign machine readable semantics to it. The standard also defines ways to represent and interchange this data between analytics in favor of interoperability and platform-independence.

¹<http://docs.oasis-open.org/uima/v1.0/uima-v1.0.html>, last accessed on 2018-09-03.

Apache UIMA is the open-source implementation of said UIMA specification. A common problem with Apache UIMA is scaling [DCR⁺15, ESI⁺12, RBB⁺10]. It provides two distinct interfaces to analyze larger collections of unstructured data, with one being UIMA-AS and the other being the more dated and less flexible CPE [FLVN09]. Apache UIMA is available for Java and C++, while its scaling solutions, UIMA-CPE and UIMA-AS are only available for Java, which is why this thesis focuses on the Java implementation. Since UIMA and Apache UIMA have very similar names, which may lead to confusion, it is common practice to call the implementation simply UIMA and explicitly state when talking about the specification. This practice will be adopted for the rest of the thesis.

2.1.1 Apache UIMA

Apache UIMA is one of few general approaches to implement NLP solutions and the only commonly known implementation of the specification with the same name. With a very modular architecture, UIMA is a popular tool that can be applied to a majority of NLP problems. A large part of the popularity of UIMA stems from the large DKPro Core collection of components, containing hundreds of analysis modules and precomputed language models [EdCG14], which can be imported into existing Java projects with the build automation tool Apache Maven [DKPb].

UIMA is usually used to process not a single but whole corpora of documents. A document in this sense is text, although the UIMA specification permits other data types as well. However, UIMA can not yet handle other data types without serializing it first. The UIMA specification, as well as the implementation do not directly pose limitations to the document size but since documents are stored in native Java String variables, which themselves are implemented as arrays of chars, the practical limit of documents sizes is dependent on the JVM version and is around one to two gigabytes [Staa] per document. In the context of UIMA, such a document is called a SofA (Subject of Analysis).

An analytic in the UIMA specification is called an *Analysis Engine* (AE) in the implementation. An AE (Analysis Engine) is code, that gets an input CAS (Common Analysis System) and produces a number of analysis results on the SofA. Common examples for AEs are Segmentation, Tokenization, and Part-of-Speech finding algorithms, which are basic NLP tasks [DKPb]. However, since an AE contains arbitrary Java code, any form of analysis can be instrumentalized by UIMA. It is defined by an XML Analysis Engine descriptor. Such an AE can either be a so-called *primitive* or an *aggregate* engine. An aggregate engine simply contains one or more other AEs, that are aggregated into one single engine.

Analysis results are stored as annotations. An annotation has at least two attributes `int begin` and `int end`, indicating the start and end index of the SofAs substring this annotation is associated with. This concept is extendable to any kind of SofA that contains any sort of subsets, for example images or audio and video streams. However, this is impractical for reasons mentioned above. It is possible, but uncommon, to define other types of subsets on a string that – for example – permit multiple segments. Such

subsets can be implemented in a custom implementation of the `AnnotationBase` class, which in turn may omit the concept of a `begin` and `end`. Since an important reason of the popularity of UIMA is due to the large AE repositories and the possibility to reuse already published code, custom annotation implementations are rarely used because it would most likely lead to incompatibilities. However, subclassing the `Annotation` class is often done to ensure type safety, leading to hierarchies. An example hierarchy can be seen in Figure 2.1. Here, a `Div` annotation inherits from the `Annotation` class and acts as a superclass for three other annotations, each describing parts of a SofA. The classes `Split` and `Compound` also directly inherit from `Annotation`, but are related with each other. A compound consists of multiple splits, which can also contain splits themselves. A `Split` annotation may be either a `CompoundPart` or a `LinkingMorpheme`.

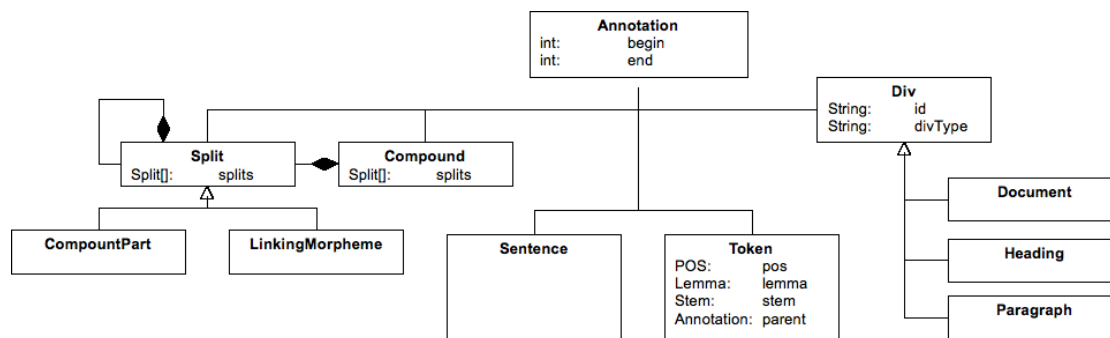


Figure 2.1: An example UIMA type hierarchy [DKPa]. Everything inherits from the `Annotation` class. Annotations may have attributes containing other annotations, or arbitrary Java objects in general.

Each annotation type must be part of an underlying Type System. The Type System is a schema of all available types of annotations that may be associated with a current SofA, thus it provides the meta data for the annotations. It is defined by an XML Type System descriptor that is usually used by the *JCasGen*, a Java code generator for UIMA types. On creation, a parent Type System can be specified, allowing inheriting from types that are not defined in the current context and encapsulate all in a single larger Type System.

The SofA, all analysis results in form of annotations that are compliant to an underlying Type System, and the Type System itself are stored together in one wrapping object, called a Common Analysis Structure (CAS). It is the sole input an AE gets, since it incorporates the complete context needed for the analysis. To store annotations efficiently, different indexes are created, providing fast access to commonly used queries like overlapping or distinct annotations and preventing deletion or modification of already established analysis results. Furthermore, a CAS object provides different Views, lightweight versions of a CAS, that store their own SofA and annotation index. These Views are identified by a String, while the original data of the CAS is usually called the *Initial View*.

A *Collection Reader* implements an interface very similar to the well-known `Iterator`, namely it provides the functions `boolean hasNext()` and `CAS getNext()`. A *Collection Reader* usually takes the role of initializing the CAS with the `SofA`. Well known *Collection Readers* achieve this by reading from a file system, a database or `Collection` object, but any other collection may be read by implementing a custom *Collection Reader*. It is also configured by writing an XML descriptor file.

Multiple Analysis Engines that form a complete flow of analysis are commonly known as a *pipeline*. Since multiple AEs can be aggregated into one, a pipeline is usually an instance of a single aggregate Analysis Engine. Sometimes a pipeline is meant to also include a *Collection Reader*, however this will not be the case in this thesis. Because of the convention to call analysis results annotations, AEs are often called *Annotators*, which is not correct in general, since engines do not need to attach any annotations to the input CAS.




Figure 2.2: An example UIMA pipeline for named entity recognition [Apac].

Figure 2.2 shows a simplified view of an analysis pipeline for named entity recognition. Given a CAS by a *Collection Reader* (omitted here), the pipeline which is really just a aggregate Analysis Engine starts to identify the language of the text with an Analysis Engine specifically designed for that task. This AE stores the results inside the CAS and forwards it to the next engine in line, which is a tokenizer.

A tokenizer annotates words, sentences and punctuation and is highly language dependent. It uses the analysis result given by the AE before deciding upon an algorithm or model to use according to the detected language. After tokenization, a Part-Of-Speech Tagger annotates each words part of speech with a different annotation according to a tag set. There are a number of tag sets for different applications, as for example the Penn Treebank Project tag set¹ and the STTS (Stuttgart-Tübingen-TagSet). The Part-Of-Speech Tagger is highly language dependent because of this. It also utilizes the results of the tokenizer, since it iterates over all annotations that are words (as opposed to annotations that span punctuation).

Afterwards the CAS gets put into a Shallow Parser, which analyzes Part-Of-Speech tags and their semantic relation among other tags in the same sentence. In a sentence ‘I like green apples.’ a Part-Of-Speech Tagger would correctly decide that ‘green’ is an adjective and ‘apples’ is a noun. However, a parser would combine those two to form

¹https://www.ling.upenn.edu/courses/Fall_2003/ling001/penn_treebank_pos.html, last accessed on 2018-09-08.

‘green apples’, a Noun Phrase, because ‘green’ is an adjectival modifier of ‘apples’. A Parser may also be used to improve the results of a previous Part-Of-Speech tagging.

A Named Entity Recognizer then takes the CAS object and looks for fitting entities. This is commonly an entry of a given noun list, but can be more sophisticated, depending on the wanted precision, the entity type and computation speed. After the last part of the pipeline returns, the analysis is done. The resulting CAS now includes a number of analysis results in form of annotations which can now be extracted or processed further.

2.1.2 UIMAfit

Since UIMA needs XML descriptor files to configure and describe most of its components, especially pipelines and type systems, developing in it requires a large amount of manual XML configuration that leads to code that is hard to maintain. Apache UIMAfit is a framework that builds on UIMA, providing an interface to programmatically describe, instantiate, and deploy UIMA components [OB09]. UIMAfit also provides an interface to dynamically write XML descriptor files for UIMA components. However, since it is able to instantiate and deploy said components without the need of XML files, those are mostly ignored. UIMA-AS, a native UIMA scaling framework described in Section 2.1.4, is known to be widely incompatible with UIMAfit, which is what led to the creation of Leo, described in Section 5.2.1.

UIMAfit has been part of the Apache UIMA project since 2012 and is therefore officially supported [dC].

2.1.3 UIMA-CPE

UIMA-CPE was the first method to add distributed computation capability to UIMA. Nowadays it has been replaced by UIMA-AS and is mostly obsolete. It made use of so called *CAS Consumers*, engines that do not analyze the CAS, but extract the needed analysis results from it and process the data as wanted. Common uses for CAS Consumers are writing analysis results into a database or serializing the whole CAS into a XMI (XML Metadata Interchange) file. CAS Consumers have been deprecated and replaced by AEs since 2006, mainly because CAS Consumers do not provide any new functionality or are semantically different from Analysis Engines. Historically a CAS Consumer would not modify a CAS object. This convention of a reading-only and consuming Analysis Engine is often used to provide maximum modularity among UIMA engines.

Another concept exclusive to UIMA-CPE are CAS Initializers, which also have been deprecated for over a decade, but are still included in UIMA. A CAS Initializer was responsible to populate a CAS from an object given by the Collection Reader. It therefore implemented the function `initializeCas(Object document, CAS cas)`. This was used for more complex collection reading capabilities, therefore CAS Initializers are generally seen as a plug-in to Collection Readers to extend their functionality. If – for example – only table of contents of larger documents are meant to be analyzed, then a Collection Reader would read the whole document and pass it to the CAS Initializer, which would

search for a table of contents and fill the CAS with its findings and discard the rest of the document.

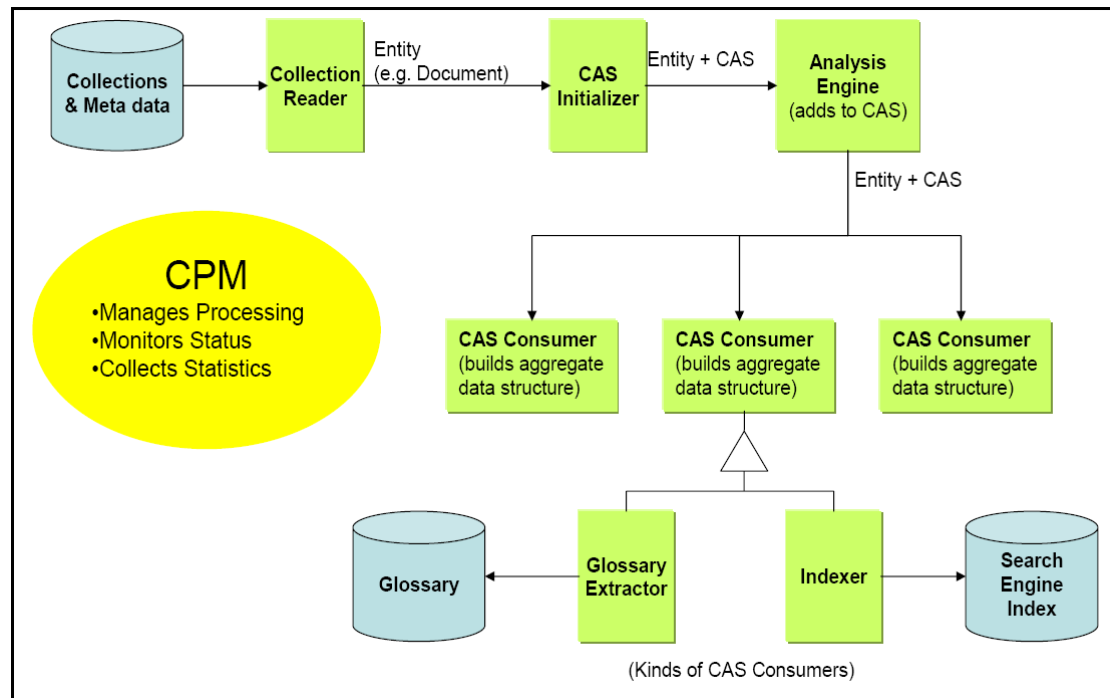


Figure 2.3: All UIMA-CPE components [Apad].

Figure 2.3 shows a complete example pipeline. It starts with any kind of collections, potentially containing meta data. A common example would be a folder hierarchy with ‘last modified’ timestamps. The Collection Reader is aware of this collection and implements an `Iterator` like interface, returning plain `Object`s. These are given to the CAS Initializer. Notice, that the CAS Initializer must be aware of what kind of entity the Collection Reader sends it. The CAS Initializer then fills a CAS object with some data from the given input `Object`. It might also create some first annotations to store meta data inside the CAS, such as the source document URL (Uniform Resource Locator) or the creation timestamp.

The CAS is then sent to the pipeline, containing one or more AEs, providing analysis results in form of annotations that are stored inside the CAS. Notice that the corresponding CAS object for a document always stays the same identical object. A CAS and its corresponding document are therefore closely associated to each other. After the analysis phase, the CAS is sent to the CAS Consumers. Those aggregate the analysis results and process it further. This process is commonly the indexing into a database or printing logs to a log file or console. Since CAS Consumers have read-only access to the CAS object they get, all of them might be processed in parallel, provided that the consumers do not interfere with each other.

All these components in combination with the UIMA CPM (Collection Processing Manager) forms the UIMA-CPE. The Collection Processing Manager provides configuration options for deployment, instantiation, and error recovery. It monitors the whole process and collects statistics. By configuration of the CPM scaling is possible either locally or on distributed machines.

For all three components introduced in this Section 2.1.3 XML descriptor files are needed for configuration. The concept of a UIMA-CPE is widely incompatible with UIMAFit, described in Section 2.1.2. UIMAFit is able to instantiate a CPE, but relies on some hardcoded default configuration, making complex multithreading applications impossible¹.

2.1.4 UIMA-AS

UIMA-AS is the successor of UIMA-CPE, providing more flexibility for scaling and deploying than its predecessor. It deploys AEs as services and registers them at a broker. UIMA-AS ships with a preconfigured instance of Apache ActiveMQ, which is an open source message broker that implements the Java Messaging Service (JMS). However, other broker implementations can also be used, but must be configured for usage with UIMA-AS and JMS. If a UIMA-AS client queries the broker, it submits a serialized CAS object to the input queue that is responsible for the wanted analysis. When any registered service finishes its current job, it pulls a new CAS from the broker and starts processing. This analysis process can also be multithreaded inside a single service. This is configurable by the deployment XML descriptor files of the AEs, but must be handled with care since multiple instances of Analysis Engines in the same JVM share static resources. After finishing the process, either successfully or by failing, the service returns the CAS object to the brokers corresponding output queue where it waits until the broker finds time to forward it to the waiting client. The described process can be seen in Figure 2.4. The user-defined AE get wrapped by a UIMA-AS controller, that handles communication with the input and output queue. These queues are provided by a broker, here ActiveMQ.

To provide capabilities of a complex flow of analyses instead of the simple synchronous order, the user can implement what is called a *Flow Controller*. An aggregate Analysis Engine can have at most one Flow Controller, that handles what AE gets the CAS next. Usually UIMA defaults to the `FixedFlow` class, which executes AEs one after another, but more sophisticated flows can be implemented. If an aggregate Analysis Engine contains such a Flow Controller, further queues are installed. Figure 2.5 shows such an advanced pipeline containing a flow controller and two delegate Analysis Engines. Submitting a CAS to the aggregate Analysis Engine queues it into the queue of the Flow Controller (FC). When it finishes its current computation, the Controller is faced with the decision which delegate Analysis Engine should obtain the CAS for processing and appends it to the corresponding queue. Notice that said queue is also provided by ActiveMQ (or

¹<https://uima.apache.org/d/uimafit-2.0.0/api/org/apache/uima/fit/cpe/CpeBuilder.html>, last accessed on 2018-09-08.

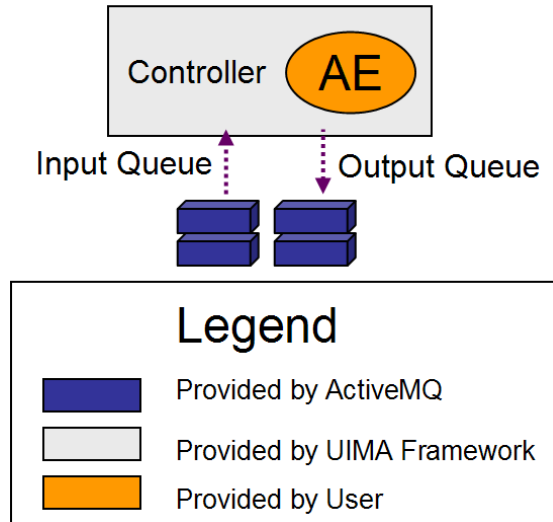


Figure 2.4: An Analysis Engine as a service in UIMA-AS [Apab].

any other implementation). After analysis, the CAS is sent back to the Flow Controller, or more specifically its output queue. It may now decide to send the processed CAS to another delegate AE or stop processing and output it to the brokers output queue. The large amount of queues may seem excessive, but it is necessary to provide a synchronous execution of the Flow Controller and – if configured – the Analysis Engines without loss of data.

Since the CAS object contains the SofA, all analysis results, the type system, and maybe even different Views, it can grow quite large over the span of a complicated pipeline. This forms a problem in UIMA-AS, since the CAS has to be serialized for every transport inside the system, from client to broker, from broker to service, from Flow Controller to delegate Analysis Engines and the whole way back. Serialization however is an I/O intensive task and even if the underlying NLP analysis contains a large amount of analysis engines, serialization might not be negligible. Epstein et al. handle this problem in [ESI⁺12] by avoiding serialization on local instances and introducing a sparse Delta-CAS serialization, containing only changes in respect to an original CAS.

As most parts of the native UIMA framework, UIMA-AS is configured by writing an XML descriptor, containing all the necessary data for deployment. A dynamic creation of said descriptor files is currently not possible with UIMAfit, but is provided by Leo, described in Section 5.2.1.

2.2 Distributed Computation

When handling large sets of data, a single machine may not be sufficient to solve the given task in a feasible time. In such a scenario, it would be desirable to just add more computation power to finish said task quicker. However, even if the given task

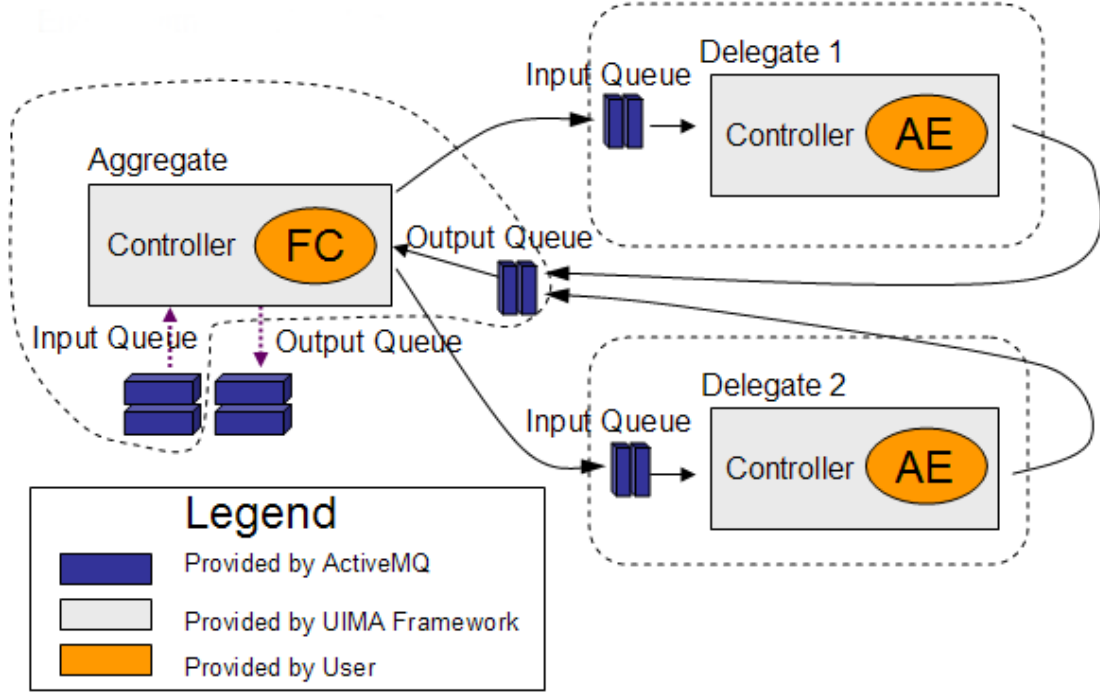


Figure 2.5: An aggregate Analysis Engine as a service in UIMA-AS [Apab].

permits parallelization, which is not obvious in general, distributing a problem among multiple machines, or similarly processing cores, is not trivial. For many problems a large administrative and communicative overhead aggravates the effort to parallelize.

In this section, some models for parallel and distributed computation are described. Those concepts aim to generalize the problem of distributing a task while at the same time try not to be too restrictive in their interface.

2.2.1 MapReduce

MapReduce is a programming model for distributed computation of large sets of data on clusters of processing cores and usually multiple machines. Google introduced the MapReduce model in 2003 and used it a few years before announcing 2014 to switch to a less restrictive framework [DG08].

The MapReduce process consists of three phases, *Map*, *Shuffle*, and *Reduce*. The shuffle phase is usually provided by an implementing framework, both other phases are to be implemented by a user. Let the input data set be C , with identifiers I . More specifically this means that the following holds:

$$\forall c \in C : \exists ! i \in I : i \text{ is associated with } c.$$

Furthermore let K be a set of keys and V a set of intermediate values. Then the `map` function maps the input data with the associated identifier to a list of key-value pairs:

$$\text{map} : I \times C \rightarrow (K \times V)^*$$

Then the `reduce` function reduces a key and a list of all associated intermediate values to a single intermediate value:

$$\text{reduce} : K \times V^* \rightarrow K \times V$$

The `reduce` function is often described with a range of just V , because it never changes its parameter of K and just passes it through. After all value lists have been reduced to contain only a single value $v \in V$, they form the final output $(K \times V)^*$.

The canonical example for this model is the problem of counting the number of occurrences for each word in large documents or even larger corpora of documents [DG08]. Recall that for given input documents C and corresponding identifiers I , which might be filenames or URLs, one expects a list of key-value pairs containing $k \in K$, an identifier for a single word (likely the word itself), and $v \in V$ an integer value describing the word's occurrences. Listing 2.1 shows an example implementation of said behavior. Notice that the pseudocode class `Word` does refer to a substring containing a single word and not the unit of data.

First, all documents C and their identifying information I are put into $|C|$ instances of the `map` function. The results are $|C|$ lists of word-integer pairs. Notice that these intermediate results are not yet distinct. This means that several entries of even the same list might be equal if the corresponding word occurs more than once in one document. This is followed by the shuffle operation, which collects intermediate results with the same key on as few machines as possible. This is a costly procedure, since data must be sent over the network. In the third phase, the reduction algorithm gets a word and a number of corresponding counting integers which it just adds and returns. Notice that – in this example – the first execution of the `reduce` function will receive a word and a list of ones. This is because the `map` function initialized each word counter with exactly one.

All intermediate results per word can now be reduced further until only one value remains, which is the final output value. Since the MapReduce model does not define an ordering on the lists given to the `reduce` function, it must be associative and commutative to always yield the same result regardless of the inputs ordering. However, MapReduce implementations usually guarantee a fixed ordering to simplify programming the `reduce` function.

Dean and Ghemawat found in [DG08] that many real world applications are describable in the MapReduce model. However, it is still very restrictive and has been abandoned by Google for this very reason. A popular open-source implementation of MapReduce is Apache Hadoop, or more specifically Hadoop MapReduce. It is therefore part of Apache Hadoop, a collection of utilities to handle large amount of data in computation. Apache Hadoop is popular for the HDFS (Hadoop Distributed File System), a high performance distributed file system.

```

1 List<Pair<Word, Integer>> map(Id docIdentifier, Text docText) {
2     List<Pair<Word, Integer>> result = new List<>();
3     for(Word w in documentText) {
4         result.append((w, 1));
5     }
6     return result;
7 }
8
9 Pair<Word, Integer> reduce(Word w, List<Integer> intermediate) {
10     Integer result = 0;
11     for(Integer i in intermediate) {
12         result += i;
13     }
14     return (w, result);
15 }

```

Listing 2.1: Example pseudocode implementation of the MapReduce model to count word occurrences.

2.2.2 Resilient Distributed Datasets

RDDs (Resilient Distributed Datasets) provide an interface that are very similar to the native Java `Collection`. They were initially developed in 2012 by Zaharia et al. in [ZCD⁺12] as a response to iterative algorithms being inefficient in current computing frameworks such as MapReduce.

An RDD can be created by either of two ways. First, stable data collections such as native Java `Collection` instances or a number of files in a file system can be initialized as RDDs. The other way of creating is a deterministic operation on an already existing RDD. These operations are called *transformations* [ZCD⁺12]. Since RDDs are immutable collections, calling such a transformation on an existing RDD is the only way of obtaining the wanted resulting RDD.

Being immutable, RDDs allow to be materialized lazily. This means that the issued transformations are executed just in time, when a materialized form of the RDD is necessary. This happens on actions like counting the number of objects in the RDD or serializing it to a file. Before executing these transformations, an acyclic graph is built to represent the necessary transformations of computing said RDD. Furthermore, RDDs are sliced into partitions, atomic pieces that never are split in the context of Spark. These partitions can then be distributed among the clusters nodes and computed whenever necessary. For the distribution of said partitions, Spark utilizes the knowledge of the issued transformation to evaluate dependencies of resulting RDDs to their predecessor. This means for example, that a `count` operation that follows a large amount of transformations does not necessarily force Spark to actually compute all these transformations. For example a transformation `crossProduct` on data sets X and Y is guaranteed to result in a

collection of size $|X| \cdot |Y|$. Materialization of $X \times Y$ is not necessary. This technique of lazy transformation provides a simple way of fault tolerance since only the RDD lineage and not the complete materialized RDD itself must be replicated among the different machines.

The only current implementation of RDDs is Apache Spark, introduced along with RDDs in 2012 [ZCD⁺12].

2.3 Containerization and Docker

Containerization describes the virtualization of applications while using the existing host kernel [Wik17]. This distinguishes containerization from using a VM (Virtual Machine). A VM simulates a completely different hardware stack than the hosts and provides a full kernel. This makes it possible to run a virtual machine with an operating system different from that the host has and provides a near perfect isolation of host and guest.

This is different from a container inside a host operating system. Since one kernel is shared, a container can only use the host's system calls and hardware. This makes containers more light-weight than VMs but also less flexible in terms of the host operating system. There are also concerns about security since VMs provide an isolation layer between host and guest which Docker does not [Tur14, JC17].

Docker is a containerization tool that was first released in 2013 by Docker, Inc. A Docker container is an instantiated *image*. An image specifies the exact user space inside the starting container and is usually defined by taking a base image and installing necessary libraries or programs for the containerized application to launch. Listing 2.2 shows a Dockerfile, an exact definition of how newly created containers of the resulting image will be configured. First, the base image is defined. Here, the official image for openJDK for Java 8 was used, a slim image containing a working openJDK installation and nothing else. Next, UIMA-AS is installed, provided that the corresponding TAR file is available at the current folder. It is first added to the image and unpacked in the same step. Then the folder's name is changed in order to omit the version number. Afterwards a directory is created to function as appendix to UIMA-AS' classpath. Every JAR and CLASS file in this folder will be loaded into an UIMA-AS instance. Utilized Analysis Engines must be available to UIMA-AS' classpath to be used.

After successful installation of UIMA-AS, some environment variables must be set and the `PATH` variable must be modified. This is possible by the `ENV` command as shown in Listing 2.2. At last, the ports that this image exposes must be defined. Since UIMA-AS transfers data via the Java Messaging Server over 61616 or HTTP over 8080, both ports must be exposed. Notice that both ports are configurable inside the UIMA-AS configuration files. However, since port remapping is easily done in the docker composition and nothing else should ever run inside the dedicated UIMA-AS container, reconfiguration of the default ports is usually discouraged.

After building the image defined in Listing 2.2, one might use it to start containers or push it to repositories. Notice that multiple instances of this image may be run simultaneously on the same machine. To enable access to every instance and allow

```
1 FROM openjdk:8
2
3 # Install UIMA-AS.
4 ADD uima-as-2.10.2-bin.tar.gz /uima-as
5 RUN mv /uima-as/apache-uima-as-2.10.2/* /uima-as
6 RUN rm -rf /uima-as/apache-uima-as-2.10.2
7 RUN mkdir /uima-as/classpath
8
9 # Set environment variables.
10 ENV UIMA_HOME=/uima-as
11 ENV PATH="/uima-as/bin:${PATH}"
12 ENV ACTIVEMQ_BASE=/active-mq
13 ENV UIMA_CLASSPATH=/uima-as/classpath
14
15 # Expose necessary ports.
16 EXPOSE 61616
17 EXPOSE 8080
```

Listing 2.2: An example Dockerfile for an UIMA-AS image.

inter-container connections, a configuration file can be utilized to create what is called a *Docker composition*. This defines how different containers may communicate with each other and how running instances depend on other containers. An important feature is the limitation of hardware resources available in composition configuration files.

Chapter 3

Scaling UIMA with Spark

In this chapter, we will first discuss the choice of Spark as a distribution technology. Afterwards the framework implementation details will be documented with a special focus on data distribution, namely serialization and compression.

3.1 Selection of Technologies

In Section 2.2, two fundamentally different computation models were introduced, namely MapReduce and RDDs. Both are generic models of how to parallelize and distribute workload among multiple processing cores. While MapReduce poses more restrictions on the underlying function, it is also more widely used than RDDs. Gopalani compared 2015 in [GA15] both methods against each other, choosing Spark as the RDD implementation and Hadoop Map Reduce as the implementation for the MapReduce model. For the example case of K -Means, he finds that Spark performs roughly 50 % better in terms of speed. While K -Means is a valid choice for an algorithm that can be used in many fields, it is not representative for all parallelizable algorithms. Since this is impossible and UIMA poses no restrictions on what classes of algorithms can be run inside an AE, the choice of whether to use RDDs or a MapReduce implementation is still non-trivial.

Many algorithms in NLP work according to language models, which are usually not only language- but also domain-dependent. General purpose models often do not suffice in terms of domain specific vocabulary and a custom model must be provided [San10]. This is where Spark can be used, because it claims to perform better on *iterative tasks* than a MapReduce approach [Wil17]. With access to a given language model being a substantial part of many NLP algorithms, RDDs seem to suit the NLP needs better than a MapReduce approach. Being the only current implementation of RDDs, the choice of distribution framework falls to Spark.

3.2 Implementation

The framework presented here consists of several classes that implement different tasks. The framework's main class `SharedUimaProcessor` delegates all work to the corresponding

classes. One complete execution of the framework, such as an analysis of one corpus of documents, contains several steps to be made. First, the framework must be instantiated. This is done by the actual user. They then order the instance of `SharedUimaProcessor` to process a pipeline according to the output of a given collection reader. To accomplish this, the framework has to read the collection, wrap documents into CAS objects and send them along with the serialized pipeline description to its workers. After the analysis part is complete, the CAS objects get sent back where they are wrapped into a `AnalysisResult` object to get access.

Figure 3.1 shows the flow of documents in a UML activity diagram. After getting read it is wrapped inside a CAS and distributed among worker nodes. There the CAS are processed and then sent back. The following sections will describe these steps in detail.

3.2.1 Initialization

The initialization of the framework consists of two parts. First, since it depends on a running Spark infrastructure, one of such must be installed. Estimating the performance of algorithms on Spark clusters is possible, but non-trivial [WK15, GA15], especially because it heavily depends on the actual code being processed. Since both, UIMA and the framework presented here provide the capability to process documents with arbitrary Java code, no assessment can be given at this point. Due to the architecture of the framework the number of usable machines is capped by the number of documents. However, since the corpus to process in a situation when utilizing a scaling framework is necessary is large, this poses no sensible limitation. Another trivial bound is a minimum number of machines, since a single machine would process all CAS faster on a native UIMA instance than a Spark cluster containing only one worker could. This is because a Spark cluster still has to administrate its only worker. The CAS has to be serialized and deserialized twice. The local UIMA instance skips this.

Given a Spark cluster, or more specific, the corresponding Java object `JavaSparkContext`, the framework itself must be instantiated. This is useful to process on multiple Spark clusters within the same JVM. The class `SharedUimaProcessor` provides a constructor

```
SharedUimaProcessor(JavaSparkContext, CompressionAlgorithm, CasSerialization, Logger)
```

While the first parameter `JavaSparkContext` was explained above as providing the necessary API to Spark for the framework to use, the others have not yet been described. The `CompressionAlgorithm` and `CasSerialization` parameters are optional and may be `null`. They are implementations of interfaces provided by the framework to specify how CAS should be serialized and compressed for network transport. This is explained further in Section 3.3. The last of the constructor's arguments is an implementation of the popular logging framework interface `org.apache.log4j.Logger`.

3.2.2 Transport

Depending on how Spark is configured, the user code is either executed directly on the master node (standalone) or on an unrelated machine that sends all necessary parameters

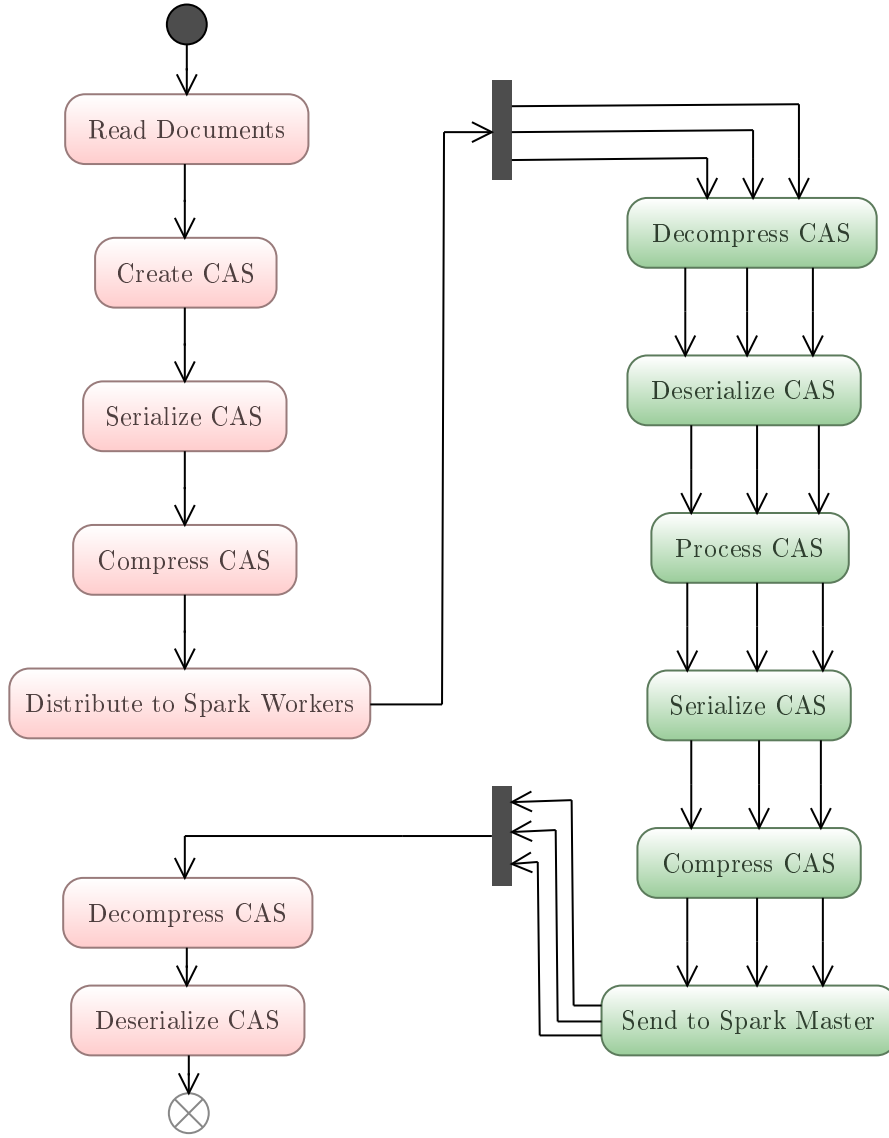


Figure 3.1: An UML (Unified Modeling Language) activity diagram of the CAS distribution process. The executing machine instantiates and distributes the CAS objects to the workers, which analyze it in parallel.

to the master node (cluster mode). Usually the standalone mode is chosen only for development or trivial clusters of a single machine, because the underlying call to execute a function is synchronous in such a configuration, therefore the process is not monitorable until the call returns. Figure 3.2 shows the whole process for a cluster mode configuration. Given the initialization described in Section 3.2.1, a collection reader would read documents into a collection of CAS objects. These are then serialized and compressed by algorithms also provided by the user. This is described further in Section 3.3. However,

after successfully compressing the CAS, it gets sent to the Spark master node. Notice that this transmission is not necessary in standalone mode, since the `SharedUimaProcessor` is then instantiated on the master node itself. Not only the CAS are needed to analyze

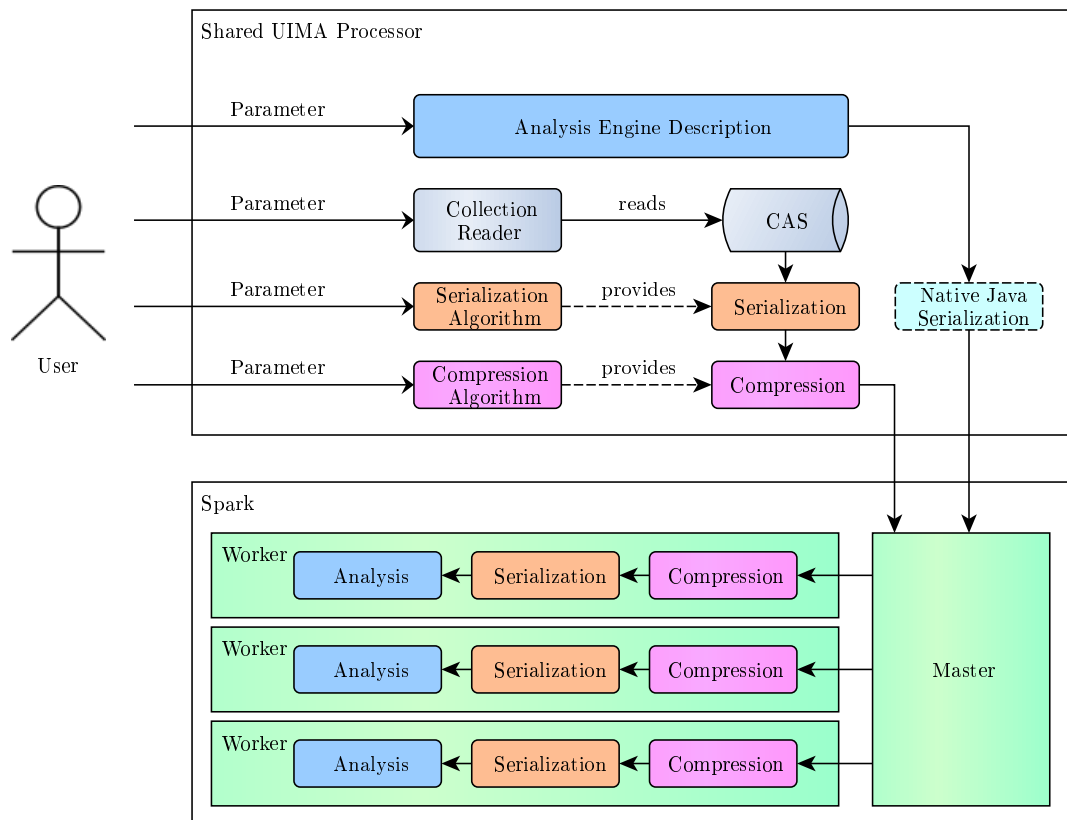


Figure 3.2: A schematic for the Shared UIMA Processor in cluster mode. The framework is instantiated with parameters, controlling most aspects of it. The CAS and the pipeline get sent to the Spark cluster, which in turn analyzes the document according to the deployed Analysis Engine.

the documents but also the analysis algorithm itself. Analysis Engines, however, are not serializable by Java, since they do not implement the required interface. This is why the framework does not accept instantiated pipelines in form of an aggregate `AnalysisEngine`, but only non-instantiated pipelines as `AnalysisEngineDescription`, which implements the `Serializable` interface. The `AnalysisEngineDescription` itself can not be executed but can be used to instantiate the corresponding AE. This happens on all worker nodes simultaneously and is combined with a non-trivial amount of computation time, since Analysis Engines may load large amount of data on the instantiation. Many NLP related algorithms need trained models or dictionaries that are relatively large [San10]. Both, UIMA

and Spark provide broadcast read-only variables to load such larger models only once, possibly saving on network and computation resources.

3.2.3 Process

As shown in Figure 3.2, one can see that a single pipeline is deployed per worker node, or more specifically per JVM. This is important to avoid a limitation of UIMAs generic nature. Since AEs consist of arbitrary code, which can not be guaranteed to be thread-safe in general. To meet the condition not to restrict any UIMA capabilities, the framework must not pose any restrictions on the Analysis Engines, which includes a guarantee for thread-safety. Instantiating exactly one pipeline per JVM circumvents the problem for the most part, as even static variables accessed by one instance are invisible to other instances. It is to mention that threading issues can still be encountered when accessing external data. The other way such problems may occur, is when deploying an aggregate Analysis Engine containing a delegate AE multiple times. In such a case a custom flow controller could be provided to execute both AEs simultaneously. However, this is also a problem in UIMAs original architecture and can be easily avoided by just using the default Flow Controller or by not adding the same Analysis Engine multiple times in one pipeline.

3.2.4 Result

The result type of the framework differs from other framework's like UIMA-AS. The resulting class, `AnalysisResult` is very similar to a `List<CAS>`, with a few but substantial differences. Figure 3.3 shows the UML class diagram of the `AnalysisResult` class. Internally it stores a `JavaRDD<SerializedCAS>`, which is a class of the Spark context. It delegates almost all commands to the underlying `JavaRDD` object, however, some functions that are sensible in the UIMA environment are also provided by this class, for example a `saveAsXmi` method, that saves all containing CAS objects into a folder. A `JavaRDD` behaves much like a `List` outside the Spark context.

The `SerializedCAS` is an internal class that represents a CAS that was serialized and compressed with the corresponding algorithms. It simply contains the serialized `Byte` array and provides an interface for deserialization by delegating the calls to the corresponding user provided algorithms. It also exposes a `size()` function to get the number of bytes needed by the compressed and serialized CAS. This is useful for evaluating algorithms that implement the `CasSerialization` and `CompressionAlgorithm` interfaces. The `SerializedCAS` class itself implements the native Java `Serializeable` interface and is therefore serializable by the JVM.

While `JavaRDD<SerializedCAS>` behaves similarly to `List<SerializedCAS>`, it is yet fundamentally different in what it does exactly. The native Java `Collection` implementations all store data on the local JVM and access them whenever they are needed. However, a `JavaRDD` is still a distributed data set among all the worker nodes that provided at least one of the resulting `SerializedCAS`. It can now be collected by the `AnalysisResult` function

`collect`. Then all CAS objects get sent back to the master node. This is a fundamental difference to UIMA-AS.

Notice that collecting all analysis results is usually *not* desired when talking about big data collections, because a single machine is likely not able to receive these large amounts of data or store it in a timely matter. Instead, a CAS Consumer should be provided at the end of the pipeline. Recall from Section 2.1.3 that a CAS Consumer is the same as an Analysis Engine in terms of implementation. However, such a CAS Consumer would extract the needed analysis results, which are most likely only a sparse subset of all given annotations, and use or store them. This storage is usually done in a database or a distributed file system like HDFS to obtain all results in one place without the need to wait for a single hard drive to write large amounts of data.

3.3 Data Distribution

Since all the input data, in form of documents, and output data, in form of analysis results, must be transmitted over a network, be it virtual or real, the serialization of larger Java objects plays a role in performance. Since both, the input and the output, are stored inside a CAS object it suffices to find a suitable serialization algorithm for those. However, finding an optimal algorithm is not trivial and usually even depends on the input data. Larger documents produce larger CAS, which in turn need a longer time to be deployed to the corresponding Spark workers. However, small documents still are no guarantee for small CAS sizes, since analysis results can be of arbitrary size and number, depending on the UIMA pipeline. Furthermore, it can be useful to compress serialized data, depending on the network setup and the serialization algorithm. Most native UIMA serializations produce XML files, which are very verbose and well compressible. Compression algorithms specifically designed for XML files achieve packing ratios of up to 80 % [GS05, MPC03, Sak09]. However, such algorithms often come at the price of a relatively high runtime. This is especially undesirable if the transmitted data is small or the serialization sparse and the expected compression ratio is low.

Since an optimal choice for both serialization and compression is not possible for the general case, the framework exposes two interfaces, namely `CasSerialization` and `CompressionAlgorithm`. Figure 3.4 shows the relationship between the framework main class `SharedUimaProcessor` and both interfaces. Additionally, two implementations that are already provided by the framework are shown in the model.

3.3.1 Serialization

In [ESI⁺12] Epstein et al. explain how serialization of CAS was an important bottleneck and a problem to solve. They configured UIMA-AS in several ways to serialize only the parts of the CAS object that are needed for further analysis. Obviously this can not be done in the general case when the underlying analysis algorithms are unknown, which is why the framework takes an instance of `CasSerialization` as an optional parameter.

```

1 public byte[] serialize(CAS cas);
2 public CAS deserialize(byte[] data, CAS cas);

```

Listing 3.1: CasSerialization method signatures

An instance of said interface implements two methods with the signatures shown in Listing 3.1.

While the signature of the `serialize` method is intuitive, this does not immediately apply to the `deserialize` function. Here, a previously created CAS object is given as a parameter for two reasons. First, UIMA allows for the configuration of a custom `CasInitializer`, which can alter the CAS object immediately after creation. Although the usage of `CasInitializers` has been deprecated since at least 2006, it is still a feature of UIMA and must therefore be taken care of [Apad]. By creating a new CAS on the target JVM, the framework first executes the `CasInitializers` and then passes the resulting CAS to the `deserialize` function. The second reason for this additional parameter is to pass the current UIMA type system. The serialized data might include annotations of types that are unknown to the native UIMA type system and therefore must be defined before deserialization. Although a parameter `TypeSystem` would have sufficed, the first reason implies the requirement of a complete CAS parameter. Since the created CAS already includes the full type system description, available by `cas.getTypeSystem()`, the framework abstains from passing another parameter to the `deserialize` method. If `CasInitializers` get removed from UIMA, this might be a feasible change in the future.

The framework already ships with two implementations of the `CasSerialization` interface, namely `XmiCasSerialization` and `UimaCasSerialization`. The `XmiCasSerialization` creates complete XMI files, containing the SofA, all analysis results and even the used type system description. To accomplish this, it uses the UIMA `XmiCasSerializer` class. Thus, the `XmiCasSerialization` implementation of `CasSerialization` acts as a mere wrapper. The second serialization algorithm `UimaCasSerialization` also just wraps around the native UIMA class `Serializer`, which is the same serialization algorithm UIMA-AS uses to distribute and retrieve CAS objects. As shown in Figure 3.4, both `XmiCasSerialization` and `UimaCasSerialization` are also implementing a singleton pattern, because no instance dependent information must be stored for either of them. However, one could implement a `CasSerialization` that stores context dependent information, for example the underlying type system.

3.3.2 Compression

Since the compression results are very dependent on the use case, data size and serialization algorithm, the framework provides the user with a `CompressionAlgorithm` interface. An implementation of said interface exposes two methods with signatures as shown in Listing 3.2. Completely abstracted from any UIMA concept, this interface simply ex-

```
1 public byte[] compress(final byte[] input);  
2 public byte[] decompress(final byte[] input);
```

Listing 3.2: CompressionAlgorithm method signatures

pects two functions, `compress` and `decompress` to behave such that for every input `byte[] X` holds:

$$X = \text{decompress}(\text{compress}(X))$$

While this is the only technical requirement for this interface, it is usually desired to have $|X| > |\text{compress}(X)|$. Since both methods act UIMA unaware, reducing the object size by omitting parts of the CAS is not possible without deserializing the CAS first, a step that is defined in the `CasSerialization` interface and not accessible from this context.

The framework ships with two implementations of the `CompressionAlgorithm` interface. It defaults to the `NoCompression` class, simply implementing the identity with $X = \text{compress}(X)$, effectively disabling any kind of compression. This is useful if network delay is negligible, especially in virtual networks inside a single machine or on low latency environments. A compression algorithm would need computation time to process all transmitted CAS, while saving only a minimum of transfer time. Secondly, the class `ZLib` implements the DEFLATE compression, which is a general purpose lossless compression algorithm, commonly used in ZIP files. As seen in Figure 3.4 both classes implement the singleton pattern, because no instance data has to be stored for either compression algorithm. However, one could implement an algorithm that stored such information, for example a complete corpus spanning dictionary.




Figure 3.3: An UML class diagram of the framework’s result classes.




Figure 3.4: An UML class diagram of the serialization and compression interfaces.

Chapter 4

Evaluation

In the following sections, the framework introduced in Chapter 3 will be compared to UIMA-AS and a single-threaded approach. For this reason, an evaluation architecture was deployed on a virtual network, which is described in Section 4.2. The comparison will follow along the metrics given in Section 1.2 ‘Implementation Specification’.

4.1 Conceptual Comparison

While the actual analysis code inside the aggregate AE is provided at the UIMA-AS services inherently, this does not hold for a Spark cluster. A Spark cluster listens for tasks until one gets provided and executes code according to a JAR file, whose URL is given in a parameter. However, the same cluster can work on other problems as well, without reconfiguration and reinitialization. Different tasks can even be processed in parallel, since the Spark is not just usable by the framework’s API.

This concept is fundamentally different from UIMA-AS, where a service is bound to exactly one aggregate Analysis Engine and therefore one pipeline. If another NLP algorithm has to be processed, another UIMA-AS service must be instantiated and registered to the broker, at another queue name. While the old services are still up, system resources for those can not be reallocated and idle until a CAS object for exactly this AE is sent to the broker.

These two concepts both have their advantages and disadvantages. The way Spark allocates resources for tasks makes it more flexible to use. Different pipelines can easily be processed without making changes to the computation cluster. A cluster of UIMA-AS services would need reconfiguration for such a case. However, since UIMA-AS services do not discard their resources after finishing the current task, a pipeline does not need to be reinitialized if another CAS should be analyzed. Since NLP algorithms commonly depend on large dictionaries or language models, this saves on loading time and especially disk I/O [San10]. A Spark cluster would immediately forget the context of the current pipeline and would have to reinitialize it.

This also comes into play when evaluating the running time of both concepts. Given a more sophisticated pipeline and a ready-to-use Spark cluster on the one hand and a

number of UIMA-AS instances on the other, the pipelines in the UIMA-AS services would have already been initialized, while the pipelines in the Spark architecture are initialized on-the-fly. This does not hold true for AEs that load their resources lazily. However, there is no guarantee for this in the general case.

4.2 Setup

Setting up a testing and benchmarking environment for the framework given in Chapter 3 and UIMA-AS is not trivial when accounting for comparability. This is because the given framework and UIMA-AS function fundamentally different from each other as described in Section 4.1. The following descriptions for Docker architectures suggest a highly similar setup, but even the initialization of both systems differ in terms of concept.

To evaluate both frameworks, and the single-threaded approach to get a sense of the administrative workload, a cluster of multiple computers was needed. To simulate a network of machines, Docker was chosen as a virtualization concept for multiple reasons. First, Dockers resource footprint is way smaller than one of a virtual machine, therefore more resources can be allocated for the actual benchmark. Second, a docker image is completely reproducible. While this can also be achieved with configuration management tools like Chef¹, Puppet² or Ansible³, defining a Dockerfile allows for easy reproduction and configuration.

Notice that network transport delay between the various components are trivial in a simulated network without any artificial delay. For this reason, compression was disabled for the framework presented in this thesis.

4.2.1 Hard- and Software

The hardware setup consisted of one Intel Xeon E5-1650 with 6x2.6 GHz with 128 GB DDR4 memory clocked at 2400 MHz. The host operating system was an Ubuntu 16.04 LTS with Docker 17.12.1-ce installed. For executing the JVMs, openJDK Docker images were used, containing the openJDK 8u181. UIMA and UIMA-AS were used in version 2.10.2 and the DKPro Core Analysis Engines are taken from the repository version 1.8.0.

4.2.2 Apache Spark

In Figure 4.1, one can see the Docker architecture used in the evaluation. The submitter reads the documents and sends the resulting CAS to the Spark master, which distributes it to its workers. The workers pull the corresponding JAR file from a source which is public to the whole cluster.

A custom created Docker image for Spark was used for both, the master and the worker nodes. This image is based on OpenJDK and will be available as described in

¹<https://www.chef.io/chef/>, last accessed on 2018-09-16.

²<https://puppet.com/de>, last accessed on 2018-09-16.

³<https://www.ansible.com/>, last accessed on 2018-09-16.

Section 5.4. The Spark worker nodes can be scaled at will by a simple Docker parameter. Both containers, *jar-provider* and *document-provider* are instances of nginx¹ HTTP servers. The document provider simply provides access to the complete corpus of documents via ordinary HTTP GET requests. These documents could have also been simply copied into the *submitter* image, however an HTTP provider server for the document corpus was chosen to equalize I/O delay, both UIMA-AS and Spark would have. This approach is also easily modifiable since the documents provided by the nginx server are a volume, pointing to a persistent folder inside the host file system. Similarly, the jar-

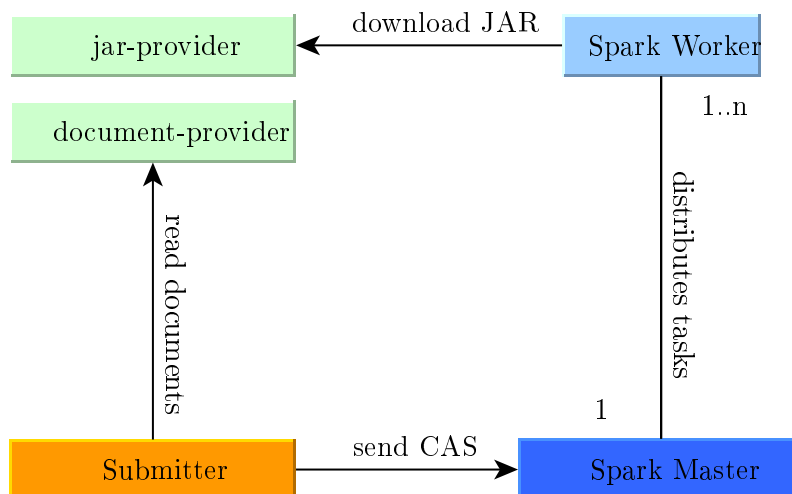


Figure 4.1: The evaluation architecture with Spark inside a Docker environment. The submitter sends the CAS objects it reads to the Spark cluster. The workers then proceed to analyze the given CAS according to code provided in a separate JAR file.

provider provides the needed JAR files for Spark worker nodes. As described above, a Spark cluster is a general computation cluster and not pipeline-specific. Thus, to execute code of an Analysis Engine, the corresponding JAR file must be provided to the whole cluster. While this can also be achieved by mounting a persistent folder inside each worker and the master node, exposing the JAR file by another nginx HTTP server is more scalable and closer to real world applications, because a JAR file will most likely not first be copied to each worker node's file system before executing, especially since it is not trivial to predict which workers are actually processing the given tasks and at what time. In this architecture, any worker can access the JAR file at any point of its lifetime.

The submitter container is also an instance of the Spark image described above, but only issues the initial command to the cluster. Spark works in one of two modes. First, the standalone mode would let the submitter container also download the JAR file and

¹<https://www.nginx.com/>, last accessed on 2018-09-16.

execute the code until a Java command issues the cluster to work on some tasks in a distributed manner. In cluster mode, a worker node would be allocated by the clusters master to execute the Java code. When a Java command orders Spark to parallelize some work, more resources are allocated for said tasks. However, the initial worker node would be unavailable for this phase. Since the architecture is more comparable to the UIMA-AS architecture in the standalone mode, it was chosen for the evaluation.

4.2.3 UIMA-AS

The Docker architecture in the evaluation setup for UIMA-AS is similar to the one for Spark. In Figure 4.2, the deployment composition is shown. For the UIMA-AS services, the underlying ApacheMQ broker and the submitter container, an UIMA-AS Docker image was composed. As with the Spark image described in Section 4.2.2, it is based on an OpenJDK image and will be available to the public according to Section 5.4. First,




Figure 4.2: The evaluation architecture with UIMA-AS inside a Docker environment. The submitter sends the CAS objects it reads to the broker, which proceeds to distribute the given tasks to instances of UIMA-AS services.

the UIMA-AS services register themselves at the given broker instance. At the same time the services initialize their corresponding pipeline. This happens in contrast to Spark, where the given cluster is completely UIMA agnostic and initializes the pipelines when they are needed.

In the same way as in Section 4.2.2, the submitter reads the documents per HTTP from the document-provider, which is an instance of the identical image used in the Spark evaluation. Notice that even the logic of reading the corpus is identical, since it is wrapped inside a `CollectionReader`, which produces CAS object that can be processed further. The CAS are then sent to the broker and distributed among the services. An

important distinction to make is that in the UIMA-AS concept the CAS object must return to the submitter. In contrast, the concept of the framework built on top of spark does allow the collection of the resulting CAS, but discourages it, since it may be a bottleneck on large corpora or large pipelines providing many analysis results.

4.2.4 Analysis Engines

Since the performance is most likely very dependent on the actual analysis to run, or more specifically the given pipeline, three different pipelines were used to compare UIMA-AS and Spark:

- Parsing
 1. *Language Setter*: Sets the CAS language to English.
 2. *Stanford Segmenter [MSB⁺14]*: Segments the text into tokens for further processing.
 3. *Stanford PoS Tagger [MSB⁺14]*: Finds the parts-of-speech for all identified tokens [TKMS03].
 4. *Malt Parser[NH05]*: Locates the grammatical components of sentences.
- Mixed Named Entity Recognition
 1. *Language Setter*: Sets the CAS language to English.
 2. *OpenNLP Segmenter [Apaa]*: Segments the text into tokens for further processing.
 3. *Mate PoS Tagger [Boh10]*: Finds the parts-of-speech for all identified tokens.
 4. *ClearNLP Lemmatizer [MSB⁺14]*: Generates the base forms for all identified tokens.
 5. *Berkeley Parser[PBTK06, PK07]*: Locates the grammatical components of sentences.
 6. *Stanford Named Entity Recognizer [MSB⁺14]*: Finds occurrences of the entities ‘Person’, ‘Location’, ‘Organization’ and ‘Misc’ and numerical entities.
- OpenNLP Named Entity Recognition
 1. *Language Setter*: Sets the CAS language to English.
 2. *OpenNLP Segmenter [Apaa]*: Segments the text into tokens for further processing.
 3. *Mate Lemmatizer [Boh10]*: Generates the base forms for all identified tokens.
 4. *OpenNLP PoS Tagger [Apaa]*: Finds the parts-of-speech for all identified tokens.
 5. *OpenNLP Named Entity Recognizer [Apaa]*: Finds occurrences of entities according to a pre-defined model.

All the given Analysis Engines (except for the Language Setter) are provided by the DKPro Core repository via the build tool Maven. All three pipelines accomplish basic NLP tasks. While the first one parses the documents grammatical components, the other two annotate named entities. However, the difference between the second and third pipeline is that the second pipeline consists of components from varying projects, while the third one contains almost exclusively OpenNLP components.

4.2.5 Document Corpus

The evaluation was used to analyze a corpus of 3036 text files, taken from the larger dataset of the project Gutenberg¹ [Lah14]. This subset was cleaned from metadata, license information and transcribers notes and were therefore seen as more realistic input data than books containing such artifacts. Furthermore books were chosen as analysis items, because they contain large amount of natural language and may revolve around very different domains. This can be important for the performance of algorithms that solve specific NLP tasks, for example Named Entity Recognition.

For the analysis, the corpus was partitioned into three slices of equal cardinality according to file size. The smallest third of all files are files up to a size of 193 kB. A file is in the middle partition, if it is larger than 193 kB, but also at most 459 kB. The largest file was about 9 MB, therefore an upper limit for even the biggest files is 10000 kB.

4.3 Results

In the following sections the results of said evaluation are described.

4.3.1 CPU Usage

Both, UIMA-AS and the framework implementation of Chapter 3 used all the available processor resources when needed. This is largely due to the fact that every (virtual) machine that represented worker nodes or services respectively, got exactly one processor core. Since no waiting time or disk I/O was necessary to process the pipelines given in Section 4.2.4, nothing stopped any of both frameworks utilizing the maximum processor time available.

Both frameworks also acted the same when idling, taking no measurable processor time. This may be due to the blocking system call `accept`² that does not introduce any CPU overhead.

4.3.2 RAM Usage

The usage of RAM was mostly neglected on the evaluation. This is because of two reasons: First, the actually used RAM is highly dependent on the underlying algorithm. It showed that both, UIMA-AS and Spark had negligible overhead relative to the memory

¹<http://www.gutenberg.org/>, last accessed on 2018-09-16.

²<http://man7.org/linux/man-pages/man2/accept.2.html>, last accessed on 2018-09-18.

actually needed by the Analysis Engines. Even simple AEs load models or static resources to process the given CAS and if the necessary resources, including the input CAS and the resulting analysis objects can not fit into the provided memory, no approach that includes UIMA is feasible. All three approaches had an actual memory overhead of less than 500 MB per machine.

The second point why RAM is neglected in this evaluation is the availability of it. The final evaluation run allocated 10 GB of memory to each processing container. This was an arbitrary choice, as 4 GB would have also sufficed, which is a low amount of memory in a cluster computation environment and one could argue that a sufficient amount of memory will always be available. Notice that an algorithm does not generally run faster with more available memory. It rather avoids swapping, using the hard drive to extend the available memory.

4.3.3 Document Throughput

Analyzing the number of documents that get processed per second shows how fast the different frameworks process the documents without taking the different document sizes in account. As one can see in Figure 4.3, the throughput for the single-threaded UIMA and the UIMA-AS implementation are similar, which seems confusing at first but will be explained in a more detailed view later. It is intuitive that larger document sizes produce smaller results at document throughput. However, Spark achieves relatively better results when given more work load in this figure. This is due to the fact that the Analysis Engine initialization is contained in the deployment of UIMA-AS, but also in the processing time of the Spark approach. As discussed earlier, this is a fundamental difference between both frameworks and can not simply be omitted, since initialization is often a costly part of an AE life cycle. However, having larger corpora of data that can be processed at once, without reinitializing the pipelines, this overhead becomes less significant and the document throughput per second rises. This can be seen as a head start for both, the single-threaded approach and UIMA-AS. This head start is sacrificed by Spark in order to provide a more flexible interface as explained in Section 4.2.

One might submit a task prematurely to the analysis benchmarks to not only load the JAR file into the Spark cluster, but also allocate resources that are loaded lazily. This approach however, fails on the architecture of Spark. While UIMA-AS and the native UIMA pipeline profit from having all resources loaded that are usually loaded lazily, Spark discards resources allocated for a specific task after completing said task. While this could be circumvented by starting a task and start measuring after the first document, Spark's architecture does not allow for this, because, as described in 2.2.2, RDDs are organized in partitions, which are atomic. This means it is not possible to only process one document, but a whole partition. This again might be circumvented by specifying a number of partitions equal to the number of elements to process, however there are two reasons against that: First, a wrong partition size affects Spark's performance negatively. Second, one does not have control over which worker node actually processes the given document and therefore can not guarantee that each worker node already processed one document. As for the similarity between the single-threaded approach and UIMA-AS,




Figure 4.3: The document throughput of each framework in the three size categories.

one can see in Figure 4.4 that both perform indeed differently, depending on the size of the corpus. Here, the size of the corpus is measured in number of files in the corresponding size category. The single-threaded instance of a native UIMA pipeline remains stagnant, which is to expect since no increase in administration overhead is necessary to process more data. However, with increasing corpus size, UIMA-AS performs better, since the broker can actually utilize the additional resources.

One can also see an increase in the performance of Spark, but it is rather slow compared to UIMA-AS. Only on large corpora of big documents, Spark achieves better results than both competitors for reasons described above.

4.3.4 Byte Throughput

Measuring the throughput in terms of bytes per second is a little more complicated than the metric before, because there are different ways to interpret the amount of bytes a process is associated with. Since pipelines usually make annotations and therefore add something to the CAS object, the output size is generally not equal to the input size. However, taking compression and serialization into account, the final output size is not equal to the size transmitted between the different components, at least in the case of Spark and UIMA-AS.

Thus, two slightly different measures are shown in Figure 4.5 and Figure 4.6. Figure 4.5 shows the number of bytes *before* processing the document and Figure 4.6 shows




Figure 4.4: The document throughput of each framework in the three size categories, detailed.

the same, but calculated with the amount of bytes the CAS object holds *after* analysis. Recall that a CAS object holds all necessary data, the SofA and all analysis results. Thus it is input and output at the same time, serving well for a size indicator.

In the first figure (Figure 4.5), one can see that the single-threaded implementation again stagnates as expected. Since different document sizes are processed exactly the same, the throughput in bytes per second does not change with increasing input sizes. This however, does not hold for UIMA-AS and Spark. One can see that UIMA-AS profits from having larger document sizes, most likely because not all available resources could be utilized on the smaller instances before they were already finished. Once every resource is used near optimally, the performance of UIMA-AS also stagnates in terms of byte throughput. Notice that this is to be expected by *any* framework with a constant amount of computation resources. Spark first drops to a lower performance when increasing the document size and then rises above both other approaches. The first drop is again most likely due to the initialization of Analysis Engines. For the same reason UIMA-AS performs better at medium input sizes, Spark starts to actually use all available resources, triggering the pipeline initialization as often as worker nodes in the Spark clusters are. Since the larger documents are bigger by an order of magnitude, this initialization again becomes more and more insignificant. Figure 4.6 shows a similar trend. However, here all algorithms perform worse at smaller input data. This might be because the CAS objects




Figure 4.5: The byte throughput of each framework in the three size categories, measured in bytes before the analysis.

do not grow linearly in size after the analysis by the given three pipelines, since fewer annotation types must be shipped within the type system. Another reason is the overhead UIMA provides itself, making processing of small data more inefficient. However, since UIMA is the underlying framework for all three approaches, every one of them suffers the consequences.

4.3.5 Maintainability

Given a working Spark cluster and the corresponding UIMA-AS service collection, changing the underlying Analysis Engine varies from each of the three evaluated approaches in multiple aspects.

AEs can be arbitrarily complex and thus, their configuration options may as well. Since each approach has to deal with this, it can be ignored while comparing all three implementations. This, however, leaves the single-threaded approach to be trivial to reconfigure. One can observe this in real life applications and stems from the very generic interface the `AnalysisEngine` class implements.

Changing AEs inside an UIMA-AS environment is substantially harder. Here, all services that serve the specific broker endpoint which is associated with the pipeline, must be shut down and redeployed. This can be automated with tools like Docker, posing other restrictions on log persistence and adding another layer of complexity to




Figure 4.6: The byte throughput of each framework in the three size categories, measured in bytes after the analysis.

the deployment logic. Notice that, without extra tooling such as Docker, each service must be shut down separately. There is no centralized command interface.

Switching AEs in a Spark environment is as simple as the single-threaded approach, since Spark is UIMA agnostic. At the moment code gets executed that demands the use of a newer version of the AE, Spark will pull the modified pipeline code and therefore process the given CAS objects within the new AE version. Notice that no changes on the cluster must be made and changing the underlying Analysis Engine is possible by only modifying the code in constant time in respect to the cluster size.

4.3.6 Implementability

Since each of the three evaluated implementations, utilizing Spark, UIMA-AS or just a single-threaded native UIMA pipeline requires knowledge about the UIMA environment, the single-threaded approach is objectively the easiest to implement. Taking sufficient knowledge of UIMA as granted, the single-threaded approach does not require any more expertise and is therefore trivial. This is not the case with UIMA-AS and the framework presented in Chapter 3.

UIMA-AS relies heavily on XML files for deployment configuration. This is an error-prone and tedious task when done by hand. The framework Leo, that builds on top of UIMA-AS helps with this by creating XML descriptor files automatically. This however,

requires the injection of the Leo framework, another dependency with possible errors which also increasing the size of the final code file. Leo is described in detail in Section 5.2.1. Having dealt with XML descriptor files and the distribution of those, other obstacles occur. Neither log files for debugging purposes nor configuration options of deployed pipeline services are available from outside the service. To make changes to a deployed pipeline, all instances of said AE must be stopped and reinitialized. Log files must be collected by hand, a custom script or by a third party tool. This is opposed to the framework built on top of Spark.

Although the framework requires knowledge about the configuration of a given Spark cluster, this has to be done only once. While UIMA-AS server instances need to be redeployed whenever a new project starts or changes on AEs occur, a Spark cluster must be set up once and can then be used indefinitely for different projects and especially distinct AEs. In an industry environment, a company-wide Spark cluster would be imaginable. In such a scenario, configuration of the Spark cluster would be done once for the whole company. Thus having knowledge about Spark is only partially a requirement, depending on the surrounding environment. Having a running and configured instance of a Spark cluster at hand, the execution of the framework becomes trivial as well. Analysis Engines can be changed whenever necessary, since Spark is completely UIMA agnostic. Log files are available by either configuring a proper log folder in the initial Spark options or by navigating to the Spark master node with any browser.

4.3.7 Code Quality

Code quality is an important factor when deciding upon a scaling framework for UIMA. Since the framework's code was optimized according to both static code analyzers FindBugs and SonarLint, it seems unfair to compare their feedback on the framework presented in this thesis with their results on analyzing alternatives like UIMA-AS or Leo. However, it is still a good indicator of code robustness.

With 3854 bad code artifacts identified by SonarLint and 42 potential bugs shown by FindBugs, UIMA-AS' code robustness is debatable. This is especially important on larger projects in a setting where frequent updates are usually not possible, for example in a production environment.

This metric has to be taken with care. As described above, the framework presented in Chapter 3 was specifically designed to output as few findings of both static code analyzers as possible. Also, both analyzers were used with the default configuration, which is usually sensible. However, if UIMA-AS would have been programmed with a different set of FindBugs and SonarLint configuration options in mind, many false positives would be expected.

Both, FindBugs and SonarLint do not find any potential bugs or vulnerabilities in the framework's code. Also the Java compiler used for compiling the framework showed no warnings.

Chapter 5

Summary

The following chapter concludes the thesis by summarizing its results and the framework's limitations. Furthermore, an outlook on the framework's source code availability and possible future changes will be given.

5.1 Limitations

The framework given in this thesis tries to not restrict any UIMA related features. To achieve this in combination with thread-safety, each pipeline is separated into their own JVM, guaranteeing maximal isolation. This however, can be unwanted, since Analysis Engines can no longer interact with each other by native Java `Thread` logic or static variables.

The `AnalysisResult` object, returned by the Spark cluster wraps around a `JavaRDD` and delegates the logic to it. However, it restricts the user from directly accessing the Spark API by setting the underlying RDD to private. This is done to isolate the user from having to handle Spark related concepts, but prevents further processing inside the cluster before collecting the data first. The framework can easily be extended to provide such functionality though, as explained in Section 5.5.

5.2 Related Work

Taking effort to scale UIMA has been done numerous times. The most prominent result was the implementation of the question answering system Watson [ESI⁺12]. This approach used native UIMA-AS, although the engineers changed the UIMA-AS source code themselves. Other approaches that are trying to be generic solutions while not posing too many restrictions or being too intrusive into the native UIMA concepts or even code, are Leo and v3NLP.

5.2.1 Leo

The Leo framework was developed by VINCI (VA Informatics and Computing Infrastructure) to allow for the easy deployment of annotators in an UIMA-AS environment [oVA]. Since it wraps around most concepts of UIMA and UIMA-AS, its architecture closely resembles UIMA-AS. This can be seen in Figure 5.1. Given a number of instances of `LeoAEDescriptor`, which are compatible with the native UIMA Analysis Engine, Leo is able to automatically write an UIMA-AS deployment descriptor file and use it to deploy a `Service` instance. This also is just a wrapper around UIMA-AS native service, which tries to register to a given broker implementation. Leo does not provide a broker implementation by its own, but depends on an existing UIMA-AS installation, which in turn provides ActiveMQ, as described further in Section 2.1.4. Given such an instance of a




Figure 5.1: The Leo architecture wrapping around UIMA-AS [oVA]

broker-service architecture, Leo is now able to perform requests to said broker. The `Leo Client` class provides access to the native UIMA-AS capability of querying the services for analysis results. For this, Leo utilizes the class `LeoCollectionReader`, which also is just a wrapper around the native UIMA Collection Readers and can be easily converted to one and vice versa.

The Leo source code repository¹ shows that the last change was in January 2018, without providing any comments or history². This, and the fact that the project only has four contributors and no pull request as of the time of writing shows that the Leo framework project is not well-maintained.

A static code analysis with FindBugs³ on the current source code⁴ shows 29 potential bugs, of which seven are rated as ‘scary’ by FindBugs. Another code analysis tool, SonarLint⁵ which also checks the coding style found a total of 626 bugs, vulnerabilities and code smells.

5.2.2 v3NLP

The framework v3NLP was first introduced in 2011 by Divita and Trietler in [DT11] as a successor to HITEx⁶. Both, HITEx and v3NLP were initially built on top of the Gate framework but have switched to UIMA since. v3NLP provides capabilities to process NLP tasks especially in the medical field. It is therefore influenced by the context of a medical environment. It provides 34 pre-composed pipelines, all related to clinical text analysis. However, it enforces the use of the CHIR⁷ Common Model, which is an ontology of labels from already existing NLP systems. This model is encoded in a default type system the framework is using and it can be extended at will. This is supposed to ensure interoperability between different NLP components that were developed for v3NLP. Notice that this poses a restriction on the user of the framework, since they are forced to always include this type system [DCT⁺16]. The framework also provides scaling capabilities that internally utilize the native UIMA-AS scalability.

While the v3NLP framework provides many functionalities for NLP research in the medical context, it is a large project in respect to UIMA, UIMA-AS and Leo (described in Section 5.2.1). v3NLP contains about 1.8 million lines of Java code in the latest commit alone as opposed to UIMA-AS with only about 140 thousand lines of code.

At the time of writing, the last commit to the framework repository was on December 2017, nine months in the past. According to the v3NLP website⁸ this is expected to be the very last commit.

5.3 Conclusion

The evaluation shows, that the framework presented here does not simply replace UIMA-AS as a method of scaling UIMA, but rather complements it. It performs generally worse

¹<https://github.com/department-of-veterans-affairs/Leo>, last accessed on 2018-09-14.

²<https://github.com/department-of-veterans-affairs/Leo/commit/038f7d5c542fa564c2997403769943ac47638692>, last accessed on 2018-09-14.

³<http://findbugs.sourceforge.net/>, last accessed on 2018-09-14.

⁴Commit 038f7d5c542fa564c2997403769943ac47638692 on 2018-09-14

⁵<https://www.sonarlint.org/>, last accessed on 2018-09-14.

⁶Health Information Text Extraction

⁷Consortium for Healthcare Informatics Research

⁸<http://inlp.bmi.utah.edu/redmine/docs/v3nlp-framework/News.html>, last accessed on 2018-09-14.

on smaller input data, be it small documents or corpora with fewer documents inside or both. However, on large input sizes it seems to be at least competitive with UIMA-AS.

The fundamental part on why it should be used as opposed to UIMA-AS is the way it is handled. Being configured once, a Spark cluster never needs to change in order to work with it. Companies or universities already owning a Spark cluster, do not need to configure it again for using it with the framework. This is different from UIMA-AS, whose Analysis Engines get deployed once and do exactly what they were designed to do.

This is why the framework presented in this thesis should be preferred on large-scale single time uses, such as an analysis of a large historic corpus, while UIMA-AS performs better on online scenarios, never having the need to reinitialize any service again. Also development and performance testing of Analysis Engines should be done with the framework given in this thesis, since the deployment of new code is orders of magnitude simpler than deploying new services with UIMA-AS. This also holds if the underlying cluster may change over time, as for example a rented cluster in the cloud, that gets increased or decreased whenever necessary.

Further evaluations with larger cluster of machines are required to find out if the trends seen in Chapter 4 also hold for industry sized Spark clusters.

5.4 Availability

From November 2018 on, the framework’s code will be publicized on GitHub¹. Since the framework is wrapped inside a Maven project, it will also be uploaded to the central Maven repository. Furthermore, another git repository that contains a working Spark dockerfile and the evaluation setup architecture will be published². Next, the Maven project that defines the benchmarking Java code will be available³. A hybrid project that consists of the deployment of UIMA-AS used in the evaluation and defines a dockerfile containing a working UIMA-AS installation will also be published on GitHub⁴.

All repositories will be published under the MIT license and are therefore free to use.

5.5 Outlook

First, the framework and all evaluation related code and resources will be made public according to Section 5.4. Other than that further improvements to the presented framework can still be made. The wrapping class `AnalysisResult` only provides a subset of its underlying `JavaRDD` functionality. The other functions were not needed at the time of writing and have therefore been neglected. However, an unwrapping of said RDD may be desired, making further processing of the underlying CAS possible. In such a way, one could benefit substantially more from Sparks optimization features.

¹On <https://github.com/s-gehring/master-thesis-program>

²On <https://github.com/s-gehring/master-thesis-spark>

³On <https://github.com/s-gehring/master-thesis-benchmark>

⁴On <https://github.com/s-gehring/master-thesis-uimaas>

Furthermore, additional compression algorithms may be implemented in the future, making compression for different serializations feasible. New serialization techniques, like delta CAS as used in [ESI⁺12] could also help improve the framework’s performance.

Apache released UIMA 3.0.0 on March 2018, which is supposed to be mostly backwards compatible to UIMA 2.10.2, for which the framework was written. Although UIMA 2.10.2 will still be supported and maintained, an upgrade to UIMA 3.0.0 also seems like a possibility for improvement.

Bibliography

- [Apaa] The Apache Software Foundation. Apache opennlp. <http://opennlp.apache.org/faq.html>. Accessed: 2018-08-29.
- [Apab] The Apache Software Foundation. Uima asynchronous scaleout. https://uima.apache.org/d/uima-as-2.8.1/uima_async_scaleout.html#ugr.async.ov. Accessed: 2018-09-09.
- [Apac] The Apache Software Foundation. Uima overview and sdk setup. https://uima.apache.org/d/uimaj-2.9.0/overview_and_setup.html. Accessed: 2018-09-08.
- [Apad] The Apache Software Foundation. Uima tutorial and developers' guides. https://uima.apache.org/d/uimaj-2.4.0/tutorials_and_users_guides.html. Accessed: 2018-02-26.
- [BL04] Steven Bird and Edward Loper. Nltk: the natural language toolkit. In *Proceedings of the ACL 2004 on Interactive poster and demonstration sessions*, page 31. Association for Computational Linguistics, 2004.
- [Boh10] Bernd Bohnet. Very high accuracy and fast dependency parsing is not a contradiction. In *Proceedings of the 23rd international conference on computational linguistics*, pages 89–97. Association for Computational Linguistics, 2010.
- [CMBT02] Hamish Cunningham, Diana Maynard, Kalina Bontcheva, and Valentin Tablan. Gate: an architecture for development of robust hlt applications. In *Proceedings of the 40th annual meeting on association for computational linguistics*, pages 168–175. Association for Computational Linguistics, 2002.
- [dC] Richard Eckart de Castilho. uimafit github repository. <https://github.com/apache/uima-uimafit>. Accessed: 2018-09-03.
- [DCR⁺15] Guy Divita, M Carter, A Redd, Q Zeng, K Gupta, B Trautner, M Samore, and A Gundlapalli. Scaling-up nlp pipelines to process large corpora of clinical notes. *Methods of information in medicine*, 54(06):548–552, 2015.

- [DCT⁺16] Guy Divita, Marjorie E Carter, Le-Thuy Tran, Doug Redd, Qing T Zeng, Scott Duvall, Matthew H Samore, and Adi V Gundlapalli. v3nlp framework: Tools to build applications for extracting concepts from clinical text. *eGEMs*, 4(3), 2016.
- [DG08] Jeffrey Dean and Sanjay Ghemawat. Mapreduce: simplified data processing on large clusters. *Communications of the ACM*, 51(1):107–113, 2008.
- [DKPa] The DKPro Core Team. Dkpro coreTM type system reference. <https://dkpro.github.io/dkpro-core/releases/1.8.0/docs/typesystem-reference.html>. Accessed: 2018-09-21.
- [DKPb] The DKPro Core Team. Dkpro coreTM user guide. [https://zoidberg.ukp.informatik.tu-darmstadt.de/jenkins/job/DKProCoreDocumentation\(GitHub\)/de.tudarmstadt.ukp.dkpro.core\\$de.tudarmstadt.ukp.dkpro.core.doc-asl/doclinks/6/user-guide.html](https://zoidberg.ukp.informatik.tu-darmstadt.de/jenkins/job/DKProCoreDocumentation(GitHub)/de.tudarmstadt.ukp.dkpro.core$de.tudarmstadt.ukp.dkpro.core.doc-asl/doclinks/6/user-guide.html). Accessed: 2018-02-26.
- [DT11] G Divita and QZ Trietler. Finding medically unexplained symptoms within va clinical documents using v3nlp. *Using administrative databases to identify cases of chronic kidney disease: a systematic review*, page 9, 2011.
- [EdCG14] Richard Eckart de Castilho and Iryna Gurevych. A broad-coverage collection of portable nlp components for building shareable analysis pipelines. In *Proceedings of the Workshop on Open Infrastructures and Analysis Frameworks for HLT*, pages 1–11, Dublin, Ireland, August 2014. Association for Computational Linguistics and Dublin City University.
- [ESI⁺12] Edward A Epstein, Marshall I Schor, BS Iyer, Adam Lally, Eric W Brown, and Jaroslaw Cwiklik. Making watson fast. *IBM Journal of Research and Development*, 56(3.4):15–1, 2012.
- [Fer12] David A Ferrucci. Introduction to “this is watson”. *IBM Journal of Research and Development*, 56(3.4):1–1, 2012.
- [FL04] David Ferrucci and Adam Lally. Uima: an architectural approach to unstructured information processing in the corporate research environment. *Natural Language Engineering*, 10(3-4):327–348, 2004.
- [FLVN09] David Ferrucci, Adam Lally, Karin Verspoor, and Eric Nyberg. Unstructured information management architecture (UIMA) version 1.0. OASIS Standard, mar 2009.
- [GA15] Satish Gopalani and Rohan Arora. Comparing apache spark and map reduce with performance analysis using k-means. *International journal of computer applications*, 113(1), 2015.

- [GS05] Marc Georges Girardot and Neelakantan Sundaresan. System and method for schema-driven compression of extensible mark-up language (xml) documents, April 19 2005. US Patent 6,883,137.
- [JC17] Zhiqiang Jian and Long Chen. A defense method against docker escape attack. In *Proceedings of the 2017 International Conference on Cryptography, Security and Privacy*, pages 142–146. ACM, 2017.
- [KSDG12] Valmeek Kudesia, Judith Strymish, Leonard D’Avolio, and Kalpana Gupta. Natural language processing to identify Foley catheter-days. *Infection control and hospital epidemiology*, 33(12):1270–1272, 2012.
- [Lah14] Shibamouli Lahiri. Complexity of Word Collocation Networks: A Preliminary Structural Analysis. In *Proceedings of the Student Research Workshop at the 14th Conference of the European Chapter of the Association for Computational Linguistics*, pages 96–105, Gothenburg, Sweden, April 2014. Association for Computational Linguistics.
- [MPC03] Jun-Ki Min, Myung-Jae Park, and Chin-Wan Chung. Xpress: A queriable compression for xml data. In *Proceedings of the 2003 ACM SIGMOD international conference on Management of data*, pages 122–133. ACM, 2003.
- [MSB⁺14] Christopher D. Manning, Mihai Surdeanu, John Bauer, Jenny Finkel, Steven J. Bethard, and David McClosky. The Stanford CoreNLP natural language processing toolkit. In *Association for Computational Linguistics (ACL) System Demonstrations*, pages 55–60, 2014.
- [NAS] NASA. Pilot study with the crew interactive mobile companion (Cimon). https://www.nasa.gov/mission_pages/station/research/experiments/2684.html. Accessed: 2018-09-17.
- [NH05] Joakim Nivre and Johan Hall. Maltparser: A language-independent system for data-driven dependency parsing. In *In Proc. of the Fourth Workshop on Treebanks and Linguistic Theories*, pages 13–95, 2005.
- [OB09] Philip Ogren and Steven Bethard. Building test suites for UIMA components. In *Proceedings of the Workshop on Software Engineering, Testing, and Quality Assurance for Natural Language Processing (SETQA-NLP 2009)*, pages 1–4, Boulder, Colorado, June 2009. Association for Computational Linguistics.
- [oVA] Departments of Veterans Affairs. Leo overview. <http://department-of-veterans-affairs.github.io/Leo/userguide.html>. Accessed: 2018-09-14.
- [PBTK06] Slav Petrov, Leon Barrett, Romain Thibaux, and Dan Klein. Learning accurate, compact, and interpretable tree annotation. In *Proceedings of the 21st*

International Conference on Computational Linguistics and the 44th annual meeting of the Association for Computational Linguistics, pages 433–440. Association for Computational Linguistics, 2006.

- [PK07] Slav Petrov and Dan Klein. Improved inference for unlexicalized parsing. In *Human Language Technologies 2007: The Conference of the North American Chapter of the Association for Computational Linguistics; Proceedings of the Main Conference*, pages 404–411, 2007.
- [PT18] Irina Pak and Phoey Lee Teh. Text segmentation techniques: A critical review. In *Innovative Computing, Optimization and Its Applications*, pages 167–181. Springer, 2018.
- [RBJB⁺10] Cartic Ramakrishnan, William A Baumgartner Jr, Judith A Blake, Gully APC Burns, K Bretonnel Cohen, Harold Drabkin, Janan Eppig, Eduard Hovy, Chun-Nan Hsu, Lawrence E Hunter, et al. Building the scientific knowledge mine (sciknowmine): a community-driven framework for text mining tools in direct service to biocuration. *Language Resources and Evaluation*, page 33, 2010.
- [Sak09] Sherif Sakr. Xml compression techniques: A survey and comparison. *Journal of Computer and System Sciences*, 75(5):303–322, 2009.
- [San10] Mark Sanderson. Manning christopher d., raghavan prabhakar, schütze hinrich, introduction to information retrieval, cambridge university press. 2008. isbn-13 978-0-521-86571-5, xxi+ 482 pages. *Natural Language Engineering*, 16(1):100–103, 2010.
- [Staa] Stackoverflow. How many characters can a java string have? <https://stackoverflow.com/questions/1179983/how-many-characters-can-a-java-string-have>. Accessed: 2018-09-03.
- [Stab] Statista. Revenues from the natural language processing (nlp) market in north america, from 2015 to 2024 (in million u.s. dollars). <https://www.statista.com/statistics/607909/north-america-natural-language-processing-market-revenues/>. Accessed: 2018-09-05.
- [Stac] Statista. Revenues from the natural language processing (nlp) market in the asia-pacific region, from 2015 to 2024 (in million u.s. dollars). <https://www.statista.com/statistics/607928/asia-pacific-natural-language-processing-market-revenues/>. Accessed: 2018-09-05.
- [Stad] Statista. Revenues from the natural language processing (nlp) market in western europe, from 2015 to 2024 (in million

u.s. dollars). <https://www.statista.com/statistics/607915/western-europe-natural-language-processing-market-revenues/>. Accessed: 2018-09-05.

- [TKMS03] Kristina Toutanova, Dan Klein, Christopher D Manning, and Yoram Singer. Feature-rich part-of-speech tagging with a cyclic dependency network. In *Proceedings of the 2003 Conference of the North American Chapter of the Association for Computational Linguistics on Human Language Technology-Volume 1*, pages 173–180. Association for Computational Linguistics, 2003.
- [TRCB13] Valentin Tablan, Ian Roberts, Hamish Cunningham, and Kalina Bontcheva. Gatecloud. net: a platform for large-scale, open-source text processing on the cloud. *Phil. Trans. R. Soc. A*, 371(1983):20120071, 2013.
- [Tur14] James Turnbull. Docker container breakout proof-of-concept exploit, 2014.
- [Wik17] Wikipedia. Operating-system-level virtualization — wikipedia, die freie enzyklopädie, 2017. [Accessed 2018-09-21].
- [Wil17] Micah Williams. Apache spark vs. mapreduce. <https://dzone.com/articles/apache-spark-introduction-and-its-comparison-to-ma>, aug 2017.
- [WK15] Kewen Wang and Mohammad Maifi Hasan Khan. Performance prediction for apache spark platform. In *High Performance Computing and Communications (HPCC), 2015 IEEE 7th International Symposium on Cyberspace Safety and Security (CSS), 2015 IEEE 12th International Conference on Embedded Software and Systems (ICESSE), 2015 IEEE 17th International Conference on*, pages 166–173. IEEE, 2015.
- [ZCD⁺12] Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J Franklin, Scott Shenker, and Ion Stoica. Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing. In *Proceedings of the 9th USENIX conference on Networked Systems Design and Implementation*, pages 2–2. USENIX Association, 2012.

Eidesstattliche Erklärung

Hiermit versichere ich, Simon Gehring, dass ich die vorliegende Masterarbeit selbstständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt habe. Die Stellen meiner Arbeit, die dem Wortlaut oder dem Sinne nach anderen Werken und Quellen, einschließlich Quellen aus dem Internet, entnommen sind, habe ich in jedem Fall unter Angabe der Quelle deutlich als Entlehnung kenntlich gemacht. Dasselbe gilt sinngemäß für Tabellen, Karten und Abbildungen.

Unterschrift: _____
Simon Gehring, Student
Universität Bonn